



# Hierarchical Bayesian modelling of cognition

**Dr. Lei Zhang**

Associate Professor

**Adaptive Learning Psychology & Neuroscience (ALPN) Lab**

Centre for Human Brain Health, School of Psychology

University of Birmingham

University of Pavia, 11-12.06.2026



[alpn-lab.github.io](https://alpn-lab.github.io)

 [@lei-zhang.bsky.social](https://bsky.social/@lei-zhang)

 [@lei\\_zhang\\_lz](https://twitter.com/lei_zhang_lz)




UNIVERSITY OF  
BIRMINGHAM




**CHBH**  
CENTRE FOR HUMAN BRAIN HEALTH

# My research:


- I ask people to **make decisions**  
**Computation**



- I build **computational models**  
**Algorithm**





- I examine **neural mechanisms**  
**Implementation**



# About me

- My journey through computational modeling
  - Started with MLE (@fminsearch in Matlab)
  - Switched to Bayesian: first JAGS, then Stan
  - Why switching to Stan? Look at Fig. 2A and Table S3 in [Zhang & Gläscher \(2020\)](#)



$$\begin{aligned}
 V_{self,t} &= [V_{self,t}(A), V_{self,t}(B)] \\
 V_{other,t} &= [V_{other,t}(A), V_{other,t}(B)] \\
 V_t &= \beta_{self} V_{self,t} + \beta_{other} V_{other,t} \\
 C1_t &\sim \text{Categorical}(\text{Softmax}(V_t)) \\
 U_{self,t} &= \beta_{self} V_{self,t} + \beta_{diff} (V_{chosen,C1,t} - V_{unchosen,C1,t}) \\
 B1_t &\sim \text{OrderedLogistic}(U_{self,t} | \theta) \\
 w.N_{opponent,t} &= \frac{\sum_{k=1}^K w_{k,t}}{\sum_{k=1}^K w_{k,t}}, K = 0, 1, \dots, 4 \\
 w.N_{with,t} &= \frac{\sum_{k=1}^K w_{k,t}}{\sum_{k=1}^K w_{k,t}} \\
 V_t(\text{switch}) &= \beta_{switch} + \beta_{diff} (V_{chosen,C1,t} - V_{unchosen,C1,t}) + \beta_{opponent} w.N_{opponent,t} \\
 C2 &\sim \text{Bernoulli}(V_t(\text{switch})) \\
 U_{self,t} &= U_{self,t} + \beta_{with} w.N_{with,t} + \beta_{opponent} w.N_{opponent,t}, \text{ if } C1 = C2 \\
 &= U_{self,t} + \beta_{with} w.N_{with,t} + \beta_{opponent} w.N_{opponent,t}, \text{ if } C1 \neq C2 \\
 B2_t &\sim \text{OrderedLogistic}(U_{self,t} | \theta) \\
 \Phi(x) &= \frac{1}{1 + e^{-x}} \\
 \delta_{self,chosen,C1,t} &= R_{self,t} - V_{self,chosen,C1,t} \\
 \delta_{self,unchosen,C1,t} &= -R_{self,t} - V_{self,unchosen,C1,t} \\
 V_{self,chosen,C1,t+1} &= V_{self,chosen,C1,t} + \alpha \delta_{self,chosen,C1,t} \\
 V_{self,unchosen,C1,t+1} &= V_{self,unchosen,C1,t} + \alpha \delta_{self,unchosen,C1,t}
 \end{aligned}$$

# Outline


- Recap on Bayesian modeling
- Why Stan?
- What is (computational) cognitive modeling?
- The idea of the Q-Learning model
- Implementing Q-Learning model for one subject in Stan
- Fitting multiple subjects with the hBayesDM package
- Q&A



# Outline

- **Recap on Bayesian modeling**
- Why Stan?
- What is (computational) cognitive modeling?
- The idea of the simple Q-Learning model
- Implementing Q-Learning model for one subject in Stan
- Fitting multiple subjects with the hBayesDM package
- Q&A

# Probability Functions






# Probability Functions

discrete variable – we talk about **mass**

- Run a test of 8 questions, and record each student's correct responses
- Count and plot the # of correct answers
- Then divide those counts by the total number of students
- The lengths of these bars sum up to 1



# Probability Functions





continuous variable – we talk about **density**

probability density function  
(PDF)

The entire area under the curve  
of PDF is 1.



cumulative distribution function  
(CDF)





# Joint, Marginal and Conditional Probability

## Joint Probability

$$p(A, B) = p(B, A)$$

- e.g.,  $p(\text{rain, cold})$ :  $p(\text{rain})$  AND  $p(\text{cold})$

## Marginal Probability

$p(A)$  – ‘ $p$  of  $A$  irrespective of  $B$ ’

- e.g.,  $p(\text{rain})$ :  $p(\text{rain, cold}) + p(\text{rain, not cold})$

## Conditional Probability

$p(A|B)$  – ‘ $p$  of  $A$  given  $B$ ’ – event  $B$  is fixed, not uncertainty

$$p(A, B) = p(A|B)p(B)$$

- e.g.,  $p(\text{rain, cold}) = p(\text{rain}|\text{cold})p(\text{cold})$

## Example: discrete

Joint probability :

$$P(X = 0, Y = 1) =$$

$$\sum_{x,y} P(X = x, Y = y) = 1$$

|      |   | snow |     |
|------|---|------|-----|
|      |   | X    |     |
| cold | Y | 1    | 0   |
|      | 1 | 0.5  | 0.1 |
|      | 0 | 0.1  | 0.3 |

Marginal probability :

$$P(Y = 1) =$$

$$P(X = 0) =$$

$$P(X = x) = \sum_y P(X = x, Y = y)$$


Conditional probability :

$$P(X = 1|Y = 1) =$$

$$\begin{aligned} P(X = x|Y = y) &= \frac{P(X = x, Y = y)}{P(Y = y)} \\ &= \frac{P(X = x, Y = y)}{\sum_x P(X = x, Y = y)} \end{aligned}$$


# Example: continuous

joint distribution




The "co-distribution" of  $x$  and  $y$ .

marginal distribution



The density of  $x$ - (or  $y$ -) values,  
without knowing the other's value.

conditional distribution



The probability distribution of  $x$ ,  
given that we know the value of  $y$ .

# Bayes' theorem

$$p(A,B) = p(B,A)$$


$$p(A,B) = p(A|B)p(B)$$

$$p(B,A) = p(B|A)p(A)$$

$$p(A|B)p(B) = p(B|A)p(A)$$

$$p(A | B) = \frac{p(B | A) p(A)}{p(B)}$$

# Linking Data and Parameter



The diagram shows two blue arrows originating from the expression  $p(A|B)$  in the equation below. One arrow points to the symbol  $\theta$ , and the other points to the symbol  $D$ .

$$p(A|B) = \frac{p(B|A)p(A)}{p(B)}$$



# Linking Data and Parameter

$$p(\theta|D) = \frac{p(D|\theta)p(\theta)}{p(D)}$$

# Linking Data and Parameter

## Likelihood

How plausible is the data given our parameter is true?

## Prior

How plausible is our parameter before observing the data?

$$p(\theta|D) = \frac{p(D|\theta)p(\theta)}{p(D)}$$

## Posterior

How plausible is our parameter given the observed data?

## Evidence

How plausible is the data under all possible parameters?


# Bernoulli Model

- You are interested in if a coin is biased.
- You will flip the coin.
- You will record whether it comes up a head (h) or a tail (t).
- You might observe 15 heads out of 20 flips.
- What is your degree of belief about the biased parameter  $\vartheta$ ?



# Bernoulli Model

$$p(h \mid \theta) = \theta^h (1 - \theta)^{1-h}$$




$\theta \sim \text{Uniform}(0, 1)$

$h \sim \text{Bernoulli}(\theta)$

|            | continuous | discrete |
|------------|------------|----------|
| unobserved | $\theta$   | $\delta$ |
| observed   | $y$        | $N$      |

# Steps of (Bayesian) Modeling?




# Solve it by Grid Approximation

discrete parameters

$$p(\theta | D) = \frac{p(D | \theta) p(\theta)}{\sum_{\theta^*} p(D | \theta^*) p(\theta^*)}$$

continuous parameters

$$p(\theta | D) = \frac{p(D | \theta) p(\theta)}{\int p(D | \theta^*) p(\theta^*) d\theta^*}$$



# What if I have multiple parameters?

grid approximation for  
2 parameters?  
5 parameters?  
10 parameters?

$$p(\theta | D) = \frac{p(D | \theta) p(\theta)}{\int p(D | \theta^*) p(\theta^*) d\theta^*}$$

$$p(data) = \int_{\text{All } \theta_1} \int_{\text{All } \theta_2} p(data, \theta_1, \theta_2) d\theta_1 d\theta_2$$

$$p(data) = \int_{\mu_1} \int_{\sigma_1} \dots \int_{\mu_{100}} \int_{\sigma_{100}} \underbrace{p(data | \mu_1, \sigma_1, \dots, \mu_{100}, \sigma_{100})}_{\text{likelihood}} \times \underbrace{p(\mu_1, \sigma_1, \dots, \mu_{100}, \sigma_{100})}_{\text{prior}} d\mu_1 d\sigma_1 \dots d\mu_{100} d\sigma_{100}$$

- Analytical solutions (often does not exist)
- Grid approximation (takes too long)
- solution: **Markov Chain Monte Carlo**

$$p(\theta | D) \propto p(D | \theta) p(\theta)$$




# MCMC Sampling Algorithms


- Rejection sampling
- Importance sampling
- Metropolis algorithm
- Gibbs sampling → JAGS
- HMC sampling\*



Stan!



# build some intuition



# Outline

- Recap on Bayesian modeling
- **Why Stan?**
- What is (computational) cognitive modeling?
- The idea of the simple Q-Learning model
- Implementing Q-Learning model for one subject in Stan
- Fitting multiple subjects with the hBayesDM package
- Q&A

# Is Stan popular?

2021 Summer

## Stan: A Probabilistic Programming Language

Carpenter, Bob; Gelman, Andrew; Hoffman, Matthew D.; Lee, Daniel; Goodrich, Ben; Betancourt, Michael; Brubaker, Marcus A.; Guo, Jiqiang; Li, Peter; Riddell, Allen

Grantee Submission, Journal of Statistical Software v76 n1 p1-32 Jan 2017

Stan is a probabilistic programming language for specifying statistical models. A Stan program imperatively defines a log probability function over parameters conditioned on specified data and constants. As of version 2.14.0, Stan provides full Bayesian inference for continuous-variable models through Markov chain Monte Carlo methods such as the No-U-Turn sampler, an adaptive form of Hamiltonian Monte Carlo sampling. Penalized maximum likelihood estimates are calculated using optimization methods such as the limited memory Broyden-Fletcher-Goldfarb-Shanno algorithm. Stan is also a platform for computing log densities and their gradients and Hessians, which can be used in alternative algorithms such as variational Bayes, expectation propagation, and marginal inference using approximate integration. To this end, Stan is set up so that the densities, gradients, and Hessians, along with intermediate quantities of the algorithm such as acceptance probabilities, are easily accessible. Stan can be called from the command line using the "cmdstan" package, through R using the "rstan" package, and through Python using the "pystan" package. All three interfaces support sampling and optimization-based inference with diagnostics and posterior analysis. "rstan" and "pystan" also provide access to log probabilities, gradients, Hessians, parameter transforms, and specialized plotting.

Descriptors: [Programming Languages](#), [Probability](#), [Bayesian Statistics](#), [Monte Carlo Methods](#), [Statistical Inference](#), [Maximum Likelihood Statistics](#), [Computation](#), [Statistical Distributions](#), [Computer Software](#)




## brms: An R package for Bayesian multilevel models using Stan

[PC Bürkner](#) - Journal of statistical software, 2017 - jstatsoft.org

... The **brms** package does not fit **models** itself but uses Stan on the back-end. Accordingly, all samplers implemented in Stan can be used to fit **brms models**. Currently, these are the static ...

☆ Save ⓘ Cite Cited by 5534 Related articles All 29 versions ⓘ

Cited by 3989




[PDF] **JAGS**: A program for analysis of Bayesian graphical models using Gibbs sampling

[M Plummer](#) - Proceedings of the 3rd international workshop on ..., 2003 - ci.tuwien.ac.at

**JAGS** is a program for Bayesian Graphical modelling which aims for compatibility with Classic BUGS. The program could eventually be developed as an R package. This article explains the motivations for this program, briefly describes the architecture and then ...

☆ ⓘ Cited by 4543 Related articles All 8 versions ⓘ

Cited by 6594




[PDF] **JAGS**: A program for analysis of Bayesian graphical models using Gibbs sampling

[M Plummer](#) - Proceedings of the 3rd international workshop on ..., 2003 - r-project.org

... **JAGS** is a program for Bayesian Graphical modelling which aims for compatibility ... **JAGS** and so avoid having to write a new program for each application. A second motivation for **JAGS** ...

☆ Save ⓘ Cite Cited by 6277 Related articles All 7 versions ⓘ

# Who are using Stan?



# Why Use Stan?



vs. BUGS / JAGS

- Less spatial correlation → effective samples
- Time to converge and per effective sample size:
  - | -  $\infty$  times faster
- Memory usage: 1–10%
- Language features
  - variable overwrite: `a = 4`, then `a = 5`
  - formal control flow (same as R)
  - full support of vectorizing & matrix calculation



Krzysztof Sakrejda  
@sakrejda

I keep getting asked why people should use [@mcmc\\_stan](#) so I wrote an answer:




"Selling" Stan  
[discourse.mc-stan.org](https://discourse.mc-stan.org)

27.03.18, 16:01


# HMC vs. Gibbs and Metropolis

Hamilton MC (HMC) implements No-U-Turn Sampler (NUTS)



- Two dimensions of highly correlated 250-dim normal
- 1,000,000 draws from Metropolis and Gibbs (thin to 1000)
- **1,000** draws from NUTS; 1000 independent draws


# Stan and Other Platforms





# Why Bayesian model estimation?


1. Not just a point estimate, but an entire **probability distribution over parameters**
2. Rather than only assuming participants are independent samples, we can model **hierarchical relationships**
3. Naturally avoid overfitting through **Bayesian Occam's Razor**, since we evaluate the model across the entire range of parameters



## But, any cons?

- HMC provides huge improvements in computational efficiency, but **mathematical foundations** are more difficult to follow, at least sometimes.
- Stan cannot sample from the posterior distribution of **discrete parameters** (e.g., [1, 2]).
  - → with additional effort, it can be achieved through marginalisation, see the Stan User Manual.
- There are a few unwritten **tips and tricks**.
  - Practise makes perfect and follow discussions on the forum.
- **Block-based** language → might seem rigid in the first place.

# A quick look at Stan



$\theta \sim \text{Uniform}(0, 1)$

$h \sim \text{Bernoulli}(\theta)$

```
data {  
  int<lower=0> N;  
  int<lower=0,upper=1> flip[N];  
}  
  
parameters {  
  real<lower=0,upper=1> theta;  
}  
  
model {  
  flip ~ bernoulli(theta);  
}
```

# Stan Language

```
data {  
  //... read in external data...  
}
```

```
transformed data {  
  //... pre-processing of data ...  
}
```

```
parameters {  
  //... parameters to be sampled by HMC ...  
}
```

```
transformed parameters {  
  //... pre-processing of parameters ...  
}
```

```
model {  
  //... statistical/cognitive model ...  
}
```


```
generated quantities {  
  //... post-processing of the model ...  
}
```

# Outline

- Recap on Bayesian modeling
- Why Stan?
- **What is (computational) cognitive modeling?**
- The idea of the simple Q-Learning model
- Implementing Q-Learning model for one subject in Stan
- Fitting multiple subjects with the hBayesDM package
- Q&A


# The idea of **computational modeling** is never new

Scientists use mathematical models to **approximate certain processes** (physical or mental), in order to explain and to predict.




# Computational modeling

Cognition as information processing



Kurt Lewin, (1936)

# Two flavours of computational modeling



Provide  
mechanistic  
explanation

Generate  
prediction/  
classification

## Construct



- Reward,
- Learning,
- Self v. Other,
- etc.

## Computational modeling

### On the Role of Theory and Modeling in Neuroscience

Daniel Levenstein,<sup>1</sup> Veronica A. Alvarez,<sup>2</sup> Asohan Amarasingham,<sup>3</sup> Habiba Azab,<sup>4</sup> Zhe S. Chen,<sup>5</sup> Richard C. Gerkin,<sup>6</sup> Andrea Hasenstaub,<sup>7</sup> Ramakrishnan Iyer,<sup>8</sup> Renaud B. Jolivet,<sup>9</sup> Sarah Marzen,<sup>10</sup> Joseph D. Monaco,<sup>11</sup> Astrid A. Prinz,<sup>12</sup> Salma Quraishi,<sup>13</sup> Fidel Santamaria,<sup>13</sup> Sabyasachi Shivkumar,<sup>14</sup> Matthew F. Singh,<sup>15</sup> Roger Traub,<sup>16</sup> Farzan Nadim,<sup>1,7\*</sup> Horacio G. Rotstein,<sup>1,7\*</sup> and A. David Redish<sup>18\*</sup>

### How Computational Modeling Can Force Theory Building in Psychological Science



Olivia Guest<sup>id</sup>, Andrea E. Martin

### Building better theories




Clare Press<sup>1, 2</sup>, Daniel Yon<sup>1</sup>, Cecilia Heyes<sup>3, 4</sup>



# Two flavours of computational modeling



# Influential perspective: Marr's 3 levels of analysis




David Marr, (1982)


# Outline

- Recap on Bayesian modeling
- Why Stan?
- What is (computational) cognitive modeling?
- **The idea of the simple Q-Learning model**
- Implementing Q-Learning model for one subject in Stan
- Fitting multiple subjects with the hBayesDM package
- Q&A

# The very short history




1998




2018

# why is it relevant?


## 1. Before conditioning




## 2. Before conditioning






## 3. During conditioning



## 4. After conditioning



# Reward Prediction error



## Predicted reward trials



## A neural substrate of prediction and reward

[W Schultz](#), [P Dayan](#), [PR Montague](#) - Science, 1997 - science.org

The capacity to predict future events permits a creature to detect, model, and manipulate the causal structure of its interactions with its environment. Behavioral experiments suggest that ...

☆ Speichern    Zitieren    Zitiert von: 12021    Ähnliche Artikel    Alle 46 Versionen



# Simple reinforcement learning: 2-armed bandit task



a simple task often used in the laboratory:

- repeated choice between  $N$  options ( $N$ -armed bandit)
- ...whose properties (reward amounts, probabilities) are learned through trial-and-error
- ...with a goal in mind: maximize the overall reward

## 2-armed bandit task





What can be your **strategies**:

1. **predict** the value of each deck
2. **choose** the best
3. **learn** from outcome to update predictions  
(repeat)




# How is prediction shaped by learning?




reward contingency 80:20



# One simple experiment: two choice task



choice presentation



action selection



outcome

what do we know?

choice & outcome

what can we measure?

choice accuracy




what do we not know?

RL update

$p(\text{choosing the better option})$

|   | subjID | trialID | choice | outcome |
|---|--------|---------|--------|---------|
| 1 | 1      | 1       | 1      | 1       |
| 2 | 1      | 2       | 1      | 1       |
| 3 | 1      | 3       | 1      | 1       |

# Q-learning Value Update



## Cognitive Model


- cognitive process
- using internal variables and free parameters

## Observation Model (Data Model)

- relate model to observed data
- has to account for noise

# Q-Learning

- The idea: **error-driven** learning
- Change in value is proportional to the difference between actual and predicted outcome



Value update:  $V_t = V_{t-1} + a * PE_{t-1}$


Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$

*Expectations on the next trial = the expectation on the current trial + learning rate \*  
prediction error (reward – current expectation)*

# Understand the learning rate

Value update:  $V_t = V_{t-1} + a * PE_{t-1}$


Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$



if  $\alpha = 0.9$

$$V_1 = 0$$

$$\begin{aligned} V_2 &= V_1 + 0.9 * (1 - V_1) \\ &= 0 + 0.9 * (1 - 0) \\ &= 0.9 \end{aligned}$$




reward contingency – 80:20

# Understand the learning rate

Value update:  $V_t = V_{t-1} + a * PE_{t-1}$

Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$



reward contingency – 80:20

# Understand the learning rate



Value update:  $V_t = V_{t-1} + a * PE_{t-1}$

Prediction error:  $PE_{t-1} = R_{t-1} - V_{t-1}$

$$V_t = (1 - a) V_{t-1} + a R_{t-1}$$

$$= (1 - a) (V_{t-2} + a (R_{t-2} - V_{t-2})) + a R_{t-1}$$


$$= (1 - a)^t V_0 + \int_{i=1}^t (1 - a)^{t-i} a R_i$$





# Is there an **optimal** learning rate?



- it hugely depends on the task!



a task = task structure + reward schedule + number of trials

An optimal learning rate exists only when the task is settled.


# Optimal learning rate?






## Choice rule: softmax



$$p_t(A) = \frac{e^{x^* V_t(A)}}{e^{x^* V_t(A)} + e^{x^* V_t(B)}}$$
$$= \frac{1}{1 + e^{-x^* (V_t(A) - V_t(B))}}$$



# Q-Learning



Value update:

$$V_{t+1} = V_t + \alpha * PE_t$$

Prediction error:

$$PE_t = R_t - V_t$$

choice rule (sigmoid /softmax):

$$p(C=a) = \frac{1}{1+e^{\tau*(v(b)-v(a))}}$$

- $\alpha$  - learning rate
- PE - reward prediction error
- V - value
- R - reward
- $\tau$  - softmax temperature

## Some terminologies

simple Q-Learning

$$V_{t+1} = V_t + \alpha (R_t - V_t)$$

$$p_t(A) = \frac{1}{1 + e^{-\tau * (v(A) - v(B))}}$$

$$LL = \text{categorical\_lpmf}(C_t | p_t)$$

simple linear regression

$$\mu = a + b X$$

$$LL = \text{normal\_lpdf}(Y | \mu, \delta)$$

# Some terminologies

simple Q-Learning

$$V_{t+1} = V_t + \alpha (R_t - V_t)$$

$$p_t(A) = \frac{1}{1 + e^{-\tau * (v(A) - v(B))}}$$

$$LL = \text{categorical\_lpmf}(C_t | p_t)$$

simple linear regression

$$\mu = a + b X$$

$$LL = \text{normal\_lpdf}(Y | \mu, \delta)$$


data / observation

free parameter / model parameter

internal/latent (model) variable

likelihood function

# How do I know which likelihood to use?







## The distribution zoo

by


Ben Lambert and Fergus Cooper

Last month: used by 285 people over 451 sessions in 41 countries

Since created: used by 4072 people over 6785 sessions in 107 countries




# Generalizing RL framework




[Palinteri et al. \(2015\)](#)


A. Trial details



[Swart et al. \(2017\)](#)




[Lockwood et al. \(2016\)](#)



[Will et al. \(2017\)](#)

**B**

|   | 👍   | 👎   |
|---|-----|-----|
| 👤 | 85% | 15% |
| 👤 | 70% | 30% |
| 👤 | 30% | 70% |
| 👤 | 15% | 85% |







# Social learning & decision-making


| Decision-making stage (value type)            | Target (and reference frame) of valuation |                                                |                                                                |
|-----------------------------------------------|-------------------------------------------|------------------------------------------------|----------------------------------------------------------------|
|                                               | Value of other people (for oneself)       | Value of other people's experiences (for them) | Value of social constellation (for normative social principle) |
| Choice (decision value and anticipated value) | Deciding whether to marry someone         | Choosing a school for your child               | Altruistic punishment of norm violations                       |
| Outcome (experienced value)                   | Being smiled at                           | Empathy with someone's pain                    | Enjoying fair distributions                                    |
| Learning (anticipated value)                  | Learning about someone's trustworthiness  | Learning about someone else's preferences      | Changing your opinion to increase social conformity            |




# learn **from** others (to benefit oneself)




Social environment




1<sup>st</sup> choice




1<sup>st</sup> bet






2<sup>nd</sup> choice



2<sup>nd</sup> bet



# learn for others (pro-social)



# Some more food for thought

What is a social RL task?

There are commonly two variations: either reward learning in social contexts (e.g. learn to expect monetary reward for a social partner) or social feedback learning (e.g. learn to expect social status or social evaluation). When the goal is to compare different types of feedback (e.g. social vs. non-social feedback), we suggest matching the feedback as closely as possible on ‘domain general’ properties, such as salience or preference

*Lockwood et al., 2016; Will et al., 2017*

Is the learning rate static across trials, or dynamically adapting along the course of the experiment?

The learning rate does not necessarily have to be constant. But in the case of the Rescorla–Wagner model (and related models), the learning rate is indeed static. A dynamic learning rate, however, is possible when other types of models are applied. Note that the interpretation of the learning rate we discussed in the main text is independent of this constant vs dynamic property

*Li et al., 2011; Mathys et al., 2011*

Is it possible to use RL models in the absence of choice data?



At least some sort of data is needed to perform model estimation. For example, skin conductance response (SCR) or pupil size response (PSR) have been used to fit RL models in associative fear learning tasks, where choice data was not available

*Li et al., 2011; Tzovara et al., 2018*

# Outline

- Recap on Bayesian modeling
- Why Stan?
- What is (computational) cognitive modeling?
- The idea of the simple Q-Learning model
- **Implementing Q-Learning model for one subject in Stan**
- Fitting multiple subjects with the hBayesDM package
- Q&A

# RL – Implementation



$$\alpha \sim \text{Uniform}(0, 1)$$



$$\tau \sim \text{Uniform}(0, 3)$$

$$p_t(C = A) = \frac{1}{1 + e^{\tau(V_t(B) - V_t(A))}}$$

$$V_{t+1}^c = V_t^c + \alpha(R_t - V_t^c)$$

Live demo of building the Q-Learning model  
model in Stan from scratch


# RL – MCMC Output






# What if we have more data?


$N = 1$





$N = 10$




# Fitting Multiple Participants




# Fitting Multiple Participants with hierarchical Bayesian analysis (HBA)



$\theta$




# Hierarchical RL Model



# Satisfied with the results?

Warning messages:

1: There were 3998 divergent transitions after warmup. Increasing adapt\_delta above 0.8 may help. See <http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup>  
2: Examine the pairs() plot to diagnose sampling problems



# Non-centered Reparameterization\*


$$\theta \sim \text{Normal}(\mu_\theta, \sigma_\theta)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$$

Stan likes **simple**  
distributions!




# Reparameterization

## Neal's Funnel

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



# Reparameterization


## Neal's Funnel

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
transformed parameters {  
  real y;  
  vector[9] x;  
  
  y = 3.0 * y_raw;  
  x = exp(y/2) * x_raw;  
}  
model {  
  y_raw ~ normal(0, 1);  
  x_raw ~ normal(0, 1);  
}
```





# Stan Sampling Parameters




| parameter               | description                        | constraint              | default |
|-------------------------|------------------------------------|-------------------------|---------|
| iterations              | number of MCMC samples (per chain) | int, $> 0$              | 2000    |
| delta: $\delta$         | target Metropolis acceptance rate  | $\delta \in [0, 1]$     | 0.80    |
| stepsize: $\varepsilon$ | initial HMC step size              | real, $\varepsilon > 0$ | 2.0     |
| max_treedepth: $L$      | maximum HMC steps per iteration    | int, $L > 0$            | 10      |

## Typical adjustments

- Increase iterations
- Increase delta
- Decrease stepsize
- Might have to increase max\_treedepth

```
funnel_fit2 <- stan("_scripts/funnel.stan",  
  iter = 4000,  
  control = list(adapt_delta = 0.999,  
    stepsize = 1.0,  
    max_treedepth = 20))
```

# Neal's Funnel: Comparing Performance

|                     | direct model                                                                        | adjusted direct model                                                                | reparameterized model                                                                |
|---------------------|-------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|
| Rhat (y)            | 1.22                                                                                | 1.1                                                                                  | 1.0                                                                                  |
| n_eff (y)           | 18                                                                                  | 42                                                                                   | 3886                                                                                 |
| runtime*            | 48.50 sec                                                                           | 50.76 sec                                                                            | 50.12 sec                                                                            |
| n_eff (y) / runtime | 0.37 / sec                                                                          | 0.82 / sec                                                                           | 77.53 / sec                                                                          |
| n_divergent         | 53                                                                                  | 0                                                                                    | 0                                                                                    |
| traceplot (y)       |  |  |  |

\*: 2 cores in parallel, including compiling time

# How about **Bounded** Parameters?



$$\begin{array}{ccc}
 \tilde{\theta} \sim \text{Normal}(0, 1) & \xrightarrow{\text{sigmoid}} & \tilde{\theta} \sim \text{Normal}(0, 1) \\
 \theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta} & & \theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \\
 \theta \in (-\infty, +\infty) & & \theta \in [0, 1]
 \end{array}$$

| constraint                      | reparameterization                                                                                                                     |
|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| $\theta \in (-\infty, +\infty)$ | $\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$                                                                               |
| $\theta \in [0, N]$             | $\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times N$                                                  |
| $\theta \in [M, N]$             | $\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times (N-M) + M$                                          |
| $\theta \in (0, +\infty)$       | $\theta = \exp(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}); \theta = \log(1 + \exp(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}))$ |

\* Probit<sup>-1</sup>: Normal cumulative distribution function (normcdf)

# Between-subject design

- each participant is tested in only one condition
- e.g., mood → IQ performance in the previous example



$$\theta_{\text{patient},i} \sim \text{Normal}(\mu_{\theta,\text{patient}}, \sigma_{\theta,\text{patient}})$$

$$\theta_{\text{control},i} \sim \text{Normal}(\mu_{\theta,\text{control}}, \sigma_{\theta,\text{control}}),$$

$$\mu_{\theta,\text{patient}} \sim \text{Normal}(\mu_{\theta}, \sigma_{\theta})$$


$$\mu_{\theta,\text{control}} \sim \text{Normal}(\mu_{\theta}, \sigma_{\theta})$$

$$\theta_{\text{patient},i} \sim \text{Normal}(\mu_{\theta,\text{patient}}, \sigma_{\theta,\text{patient}})$$

$$\theta_{\text{control},i} \sim \text{Normal}(\mu_{\theta,\text{control}}, \sigma_{\theta,\text{control}})$$


# Within-subject design

- each participant is tested under all conditions
- e.g., learning performance at the beginning vs. at the end of the semester



## Effect coding

$$\theta_2 = \theta_1 + \Delta\theta$$



## Effect coding

|            | $\theta_1$ | $\theta_2 - \theta_1$ |
|------------|------------|-----------------------|
| $\theta_1$ | 1          | 0                     |
| $\theta_2$ | 1          | 1                     |



## Effect coding

|            | $\theta_1$ | $\theta_2 - \theta_1$ | $\theta_3 - \theta_1$ |
|------------|------------|-----------------------|-----------------------|
| $\theta_1$ | 1          | 0                     | 0                     |
| $\theta_2$ | 1          | 1                     | 0                     |
| $\theta_3$ | 1          | 0                     | 1                     |

## Effect coding

$$\theta_i \sim \text{Normal}(\mu, \Sigma)$$

$$\mu = (\mu_1, \mu_2, \mu_3) = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} (\mu_1, \mu_2 - \mu_1, \mu_3 - \mu_1)^T$$

$$\Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{12} & \sigma_2^2 & \sigma_{23} \\ \sigma_{13} & \sigma_{23} & \sigma_3^2 \end{pmatrix} = \begin{pmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \end{pmatrix} \Omega \begin{pmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \end{pmatrix}$$

## Multiple IV: Factorial design

|                       |                           | FACTOR A<br>(TV VIOLENCE) |                              |
|-----------------------|---------------------------|---------------------------|------------------------------|
|                       |                           | Level 1 (A1)<br>(Violent) | Level 2 (A2)<br>(Nonviolent) |
| FACTOR B<br>(REALISM) | Level 1 (B1)<br>(Real)    | A1B1<br>(Real/Violent)    | A2B1<br>(Real/Nonviolent)    |
|                       | Level 2 (B2)<br>(Cartoon) | A1B2<br>(Cartoon/Violent) | A2B2<br>(Cartoon/Nonviolent) |

## Experimental Procedures

1. **Screening:** Online Questionnaire

2. **Baseline session**   
Reward tasks

3. **Iso sessions** (counterbalanced)

**Iso total**





Reward tasks

**Iso w media**



Reward tasks

## Reward tasks: EBDM task, RL task (counterbalanced)



|            | baseline | iso w media | iso total | baseline | iso w media | iso total |
|------------|----------|-------------|-----------|----------|-------------|-----------|
| non-social | 1        | 0           | 0         | 0        | 0           | 0         |
|            | 1        | 1           | 0         | 0        | 0           | 0         |
|            | 1        | 0           | 1         | 0        | 0           | 0         |
| social     | 1        | 0           | 0         | 1        | 0           | 0         |
|            | 1        | 0           | 0         | 1        | 1           | 0         |
|            | 1        | 0           | 0         | 1        | 0           | 1         |

# Information Criteria

AIC – Akaike information criterion


DIC – Deviance Information Criterion

WAIC – Widely Applicable Information Criterion

(Watanabe–Akaike information criterion)

finding the model that has  
the highest out-of-sample  
predictive accuracy

approximation to LOO



WAIC: using **entire posterior** distribution

AIC/DIC: using **point** estimation

# Compute WAIC from Likelihood

$$\text{WAIC} = -2 \widehat{\text{elpd}}_{\text{waic}}$$

expected log pointwise predictive density

$$\widehat{\text{elpd}}_{\text{waic}} = \widehat{\text{lpd}} - \widehat{p}_{\text{waic}}$$

$\widehat{\text{lpd}}$  = computed log pointwise predictive density

$$= \sum_{i=1}^n \log \left( \frac{1}{S} \sum_{s=1}^S p(y_i | \theta^s) \right).$$

```
lpd <- log(colMeans(exp(log_lik)))
```

$\widehat{p}_{\text{waic}}$  = estimated effective number of parameters

$$= \sum_{i=1}^n V_{s=1}^S (\log p(y_i | \theta^s))$$

```
p_waic <- colVars(log_lik)
```

# The {loo} Package

```
> library(loo)
> LL1      <- extract_log_lik(stanfit)
> loo1     <- loo(LL1)      # PSIS leave-one-out
> waic1    <- waic(LL1)    # WAIC
```

Computed from 4000 by 20 log-likelihood matrix

|          | Estimate | SE  |
|----------|----------|-----|
| elpd_loo | -29.5    | 3.3 |
| p_loo    | 2.7      | 1.0 |
| looic    | 58.9     | 6.7 |

Pareto Smoothed Importance Sampling






# Recording the Log-Likelihood in Stan

```
generated quantities {  
  ...  
  real log_lik[nSubjects];  
  ...  
  
  { # local section, this saves time and space  
    for (s in 1:nSubjects) {  
      vector[2] v;  
      real pe;  
  
      log_lik[s] = 0;  
      v = initV;  
  
      for (t in 1:nTrials) {  
        log_lik[s] = log_lik[s] + categorical_logit_lpmf(choice[s,t] | tau[s] * v);  
  
        pe = reward[s,t] - v[choice[s,t]];  
        v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
      }  
    }  
  }  
}
```

# HBA sounds good, but...



$$\mu_\alpha \sim \text{Uniform}(0,1)$$

$$\sigma_\alpha \sim \text{halfCauchy}(0,1)$$

$$\mu_\tau \sim \text{Uniform}(0,3)$$


$$\sigma_\tau \sim \text{halfCauchy}(0,3)$$

$$\alpha_i \sim \text{Normal}(\mu_\alpha, \sigma_\alpha) \tau(0,1)$$

$$\tau_i \sim \text{Normal}(\mu_\tau, \sigma_\tau) \tau(0,3)$$

$$p_{i,t}(C=A) = \frac{1}{1 + e^{\tau_i(V_{i,t}(B) - V_{i,t}(A))}}$$

$$V_{i,t+1}^C = V_{i,t}^C + \alpha_i(R_{i,t} - V_{i,t}^C)$$



$$V_{self,t} = [V_{self,t}(A), V_{self,t}(B)]$$

$$V_{other,t} = [V_{other,t}(A), V_{other,t}(B)]$$

$$V_t = \beta_{self} V_{self,t} + \beta_{other} V_{other,t}$$

$$C1_t \sim \text{Categorical}(\text{Softmax}(V_t))$$

$$U_{bet1,t} = \beta_{bias,t} + \beta_{diff,t}(V_{chosen,C1,t} - V_{unchosen,C1,t})$$

$$B1_t \sim \text{OrderedLogistic}(U_{bet1,t} | \theta)$$

$$w.N_{against,t} = \frac{\sum_{s=1}^K w_{s,t}}{\sum_{s=1}^4 w_{s,t}}, K=0,1,...,4$$

$$w.N_{with,t} = \frac{\sum_{s=1}^K w_{s,t}}{\sum_{s=1}^4 w_{s,t}}$$

$$V_t(\text{switch}) = \beta_{bias,t} + \beta_{diff,t}(V_{chosen,C1,t} - V_{unchosen,C1,t}) + \beta_{against} w.N_{against,t}$$

$$C2 \sim \text{Bernoulli}(V_t(\text{switch}))$$

$$U_{bet2,t} = \begin{cases} U_{bet1,t} + \beta_{with} w.N_{with,t} + \beta_{against} w.N_{against,t}, & \text{if } C1 = C2 \\ U_{bet1,t} + \beta_{with} w.N_{with,t} + \beta_{against} w.N_{against,t}, & \text{if } C1 \neq C2 \end{cases}$$

$$B2_t \sim \text{OrderedLogistic}(U_{bet2,t} | \theta)$$

$$\Phi(x) = \frac{1}{1 + e^{-x}}$$

$$\delta_{self,chosen,C2,t} = R_{self,t} - V_{self,chosen,C2,t}$$

$$\delta_{self,unchosen,C2,t} = -R_{self,t} - V_{self,unchosen,C2,t}$$

$$V_{self,chosen,C2,t+1} = V_{self,chosen,C2,t} + \alpha \delta_{self,chosen,C2,t}$$

$$V_{self,unchosen,C2,t+1} = V_{self,unchosen,C2,t} + \alpha \delta_{self,unchosen,C2,t}$$


# hBayesDM package

## hBayesDM




repo status **Active** build **passing** CRAN **1.0.2 – 2019-11-13** downloads **33K**

DOI **10.1162/CPSY\_a\_00002**

**hBayesDM** (hierarchical Bayesian modeling of Decision-Making tasks) is a user-friendly package that offers hierarchical Bayesian analysis of various computational models on an array of decision-making tasks. hBayesDM uses [Stan](#) for Bayesian inference.



Cited by 115



Bayesian Statistics and Hierarchical Bayesian Modeling for Psychological Science

Lecture 01




Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)  
Department of Cognition, Emotion, and Methods in Psychology


[lei.zhang@univie.ac.at](mailto:lei.zhang@univie.ac.at)  
[lei-zhang.net](mailto:lei-zhang.net)  
[https://github.com/lei-zhang/BayesCog\\_Wien](https://github.com/lei-zhang/BayesCog_Wien)

universität wien  
Fakultät für Psychologie

# Posterior predictive check: *can the model resonate with data?*





# Parameter recovery: *are parameters identifiable?*



# Model recovery: *are models identifiable?*

generative  
process

fitting  
process








# hBayesDM package

## hBayesDM

repo status **Active** build **passing** CRAN **1.0.2 – 2019-11-13** downloads **33K**

DOI **10.1162/CPSY\_a\_00002**

**hBayesDM** (hierarchical Bayesian modeling of Decision-Making tasks) is a user-friendly package that offers hierarchical Bayesian analysis of various computational models on an array of decision-making tasks. hBayesDM uses **Stan** for Bayesian inference.



Bayesian Statistics and Hierarchical Bayesian Modeling for Psychological Science

Lecture 01

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)  
Department of Cognition, Emotion, and Methods in Psychology

[lei.zhang@univie.ac.at](mailto:lei.zhang@univie.ac.at)  
[lei-zhang.net](mailto:lei-zhang.net)  
[https://github.com/lei-zhang/BayesCog\\_Wien](https://github.com/lei-zhang/BayesCog_Wien)

universität wien  
Fakultät für Psychologie

001 / 134:07




Adaptive Learning Psychology and Neuroscience





Welcome to BayesCog!
 

- Course overview
- Workshop 1: R Basics >
- Workshop 2: Probability and an introduction to Bayes' theorem >
- Workshop 3: Building simple models conceptually >
- Workshop 4: Introduction to building models in Stan >
- Workshop 5: Bernoulli and linear regression models >
- Workshop 6: Reinforcement learning models >
- Workshop 7: Hierarchical Bayesian modeling >
- Workshop 8: Model comparison >
- Workshop 9: Debugging in Stan >
- Bonus workshop: Introduction to model-based fMRI >
- Acknowledgements
- Resources



Visitors 1 / 119 contributors 2 last commit july repo size 360.2 MiB

## Welcome

This website is an adapted version of teaching materials originally made for the [award winning\\*](#) **BayesCog** seminar at the [Faculty of Psychology, University of Vienna](#), as part of the Advanced Seminar for master's students (Mind and Brain track; recorded during Summer Term 2020/2021). Further content from the [BayesCog workshop at UKE Hamburg](#) (2023) have also been added<sup>1</sup>.

**Recording:** Recordings from the original version of the course are available on [YouTube](#) (also see below). The most recent recording from the 2021 summer semester is also available on [YouTube](#).

**Outreach:** [Twitter thread](#) (being **liked 700+ times** on Twitter) summarizing the contents of the course.

**Award/Recognition:** The original course received a commendation award from the [Society for the Improvement of Psychological Science \(SIPS\)](#) (also see a [tweet](#)), as well as an ECR Teaching Award from the [Faculty of Psychology, University of Vienna](#).



<sup>1</sup> BayesCog: A freely available course in Bayesian statistics and hierarchical Bayesian modeling for psychological science

<sup>4</sup> **Lei Zhang** <sup>1,2,3</sup> and **Aamir Sohail** <sup>1,2</sup>

<sup>5</sup> **1** Centre for Human Brain Health, School of Psychology, University of Birmingham, Birmingham, UK **2** Institute for Mental Health, School of Psychology, University of Birmingham, Birmingham, UK **3** Centre for Developmental Science, School of Psychology, University of Birmingham, Birmingham, UK

<https://alpn-lab.github.io/BayesCog/>



# Recommended reading: tutorial

## Using reinforcement learning models in social neuroscience: frameworks, pitfalls and suggestions of best practices

Lei Zhang , Lukas Lengersdorff, Nace Mikus, Jan Gläscher, Claus Lamm  Author Notes

*Social Cognitive and Affective Neuroscience*, Volume 15, Issue 6, June 2020, Pages 695–707, <https://doi.org/10.1093/scan/nsaa089>

<https://doi.org/10.1093/scan/nsaa089>

ACCEPTED MANUSCRIPT

## Computational modelling of social cognition and behaviour—a reinforcement learning primer

Patricia L Lockwood , Miriam Klein-Flügge 

*Social Cognitive and Affective Neuroscience*, nsaa040, <https://doi.org/10.1093/scan/nsaa040>

Published: 30 March 2020 [Article history](#) 

<https://doi.org/10.1093/scan/nsaa040>

## Ten simple rules for the computational modeling of behavioral data



Robert C Wilson , Anne GE Collins 

University of Arizona, United States; University of California, Berkeley, United States

<https://elifesciences.org/articles/49547>

## The Importance of Falsification in Computational Cognitive Modeling

Stefano Palminteri,<sup>1,2,\*,‡</sup> Valentin Wyart,<sup>1,2,\*,‡</sup> and Etienne Koechlin<sup>1,2,\*</sup>

<https://doi.org/10.1016/j.tics.2017.03.011>

# Recommended reading: empirical work

## Catecholaminergic challenge uncovers distinct Pavlovian and instrumental mechanisms of motivated (in)action



Jennifer C Swart<sup>✉</sup>, Monja I Froböse, Jennifer L Cook, Dirk EM Geurts, Michael J Frank, Roshan Cools, Hanneke EM den Ouden<sup>✉</sup>  
Radboud University, The Netherlands; University of Birmingham, United Kingdom; Radboud University Medical Center, The Netherlands; Linguistic and Psychological Sciences, Brown University, United States; Brown University, United States

<https://elifesciences.org/articles/22169>

## COGNITIVE NEUROSCIENCE

## A brain network supporting social influences in human decision-making

Lei Zhang<sup>1,2\*</sup> and Jan Gläscher<sup>1\*†</sup>

<https://advances.sciencemag.org/content/6/34/eabb4159>

## Social threat learning transfers to decision making in humans

Björn Lindström<sup>a,b,c,1</sup>, Armita Golkar<sup>c,d</sup>, Simon Jangard<sup>c</sup>, Philippe N. Tobler<sup>b</sup>, and Andreas Olsson<sup>c</sup>

<sup>a</sup>Department of Social Psychology, University of Amsterdam, 1018 WT Amsterdam, The Netherlands; <sup>b</sup>Laboratory for Social and Neural Systems Research, Department of Economics, University of Zürich, 8001 Zürich, Switzerland; <sup>c</sup>Section for Psychology, Department of Clinical Neuroscience, Karolinska Institutet, 171 77 Stockholm, Sweden; and <sup>d</sup>Department of Clinical Psychology, University of Amsterdam, 1018 WT Amsterdam, The Netherlands

<https://www.pnas.org/content/116/10/4732.abstract>

## Modeling flexible behavior in childhood to adulthood shows age-dependent learning mechanisms and less optimal learning in autism in each age group

Daisy Crawley<sup>1†\*</sup>, Lei Zhang<sup>2,3,4†</sup>, Emily J. H. Jones<sup>5</sup>, Jumana Ahmad<sup>1,6</sup>, Bethany Oakley<sup>1</sup>, Antonia San José Cáceres<sup>1,7</sup>, Tony Charman<sup>8,9</sup>, Jan K. Buitelaar<sup>10,11,12</sup>, Declan G. M. Murphy<sup>1,9,13</sup>, Christopher Chatham<sup>4</sup>, Hanneke den Ouden<sup>10†</sup>, Eva Loth<sup>1,13†</sup>, the EU-AIMS LEAP group<sup>†</sup>

<https://journals.plos.org/plosbiology/article?id=10.1371/journal.pbio.3000908>

# Recommended reading: new direction

nature

[Explore content](#) ▾ [About the journal](#) ▾ [Publish with us](#) ▾

[nature](#) > [articles](#) > article

Article | [Open access](#) | Published: 02 July 2025

## Discovering cognitive strategies with tiny recurrent neural networks

[Li Ji-An](#), [Marcus K. Benna](#) & [Marcelo G. Mattar](#) 

[Nature](#) (2025) | [Cite this article](#)

**34k** Accesses | **354** Altmetric | [Metrics](#)

<https://www.nature.com/articles/s41586-025-09142-4>

nature

[Explore content](#) ▾ [About the journal](#) ▾ [Publish with us](#) ▾

[nature](#) > [articles](#) > article

Article | [Open access](#) | Published: 02 July 2025

## A foundation model to predict and capture human cognition




[Marcel Binz](#) , [Elif Akata](#), [Matthias Bethge](#), [Franziska Brändle](#), [Fred Callaway](#), [Julian Coda-Forno](#), [Peter Dayan](#), [Can Demircan](#), [Maria K. Eckstein](#), [Noémi Éltető](#), [Thomas L. Griffiths](#), [Susanne Haridi](#), [Akshay K. Jagadish](#), [Li Ji-An](#), [Alexander Kipnis](#), [Sreejan Kumar](#), [Tobias Ludwig](#), [Marvin Mathony](#), [Marcelo Mattar](#), [Alireza Modirshanechi](#), [Surabhi S. Nath](#), [Joshua C. Peterson](#), [Milena Rmus](#), [Evan M. Russek](#), ... [Eric Schulz](#)  [Show authors](#)

[Nature](#) (2025) | [Cite this article](#)

**148k** Accesses | **1** Citations | **633** Altmetric | [Metrics](#)

<https://www.nature.com/articles/s41586-025-09215-4>

# Recommended reading: book



# Want to learn more about Stan?

## Workshops

- [StanCon 2019 Hierarchical Models](#)
- [PyData NYC 2019](#)
- [StanCon 2018 Intro Stan](#)
- [StanCon 2018 Hierarchical Models](#)

<https://mc-stan.org/workshops/>

## Stan forum

📌 StanCon 2020. August 11-14 at Oregon State University

■ Announcements ■ stancon

StanCon 2020 will be at Oregon State University! There will be two days of tutorials followed by two days of talks, open discussions, and statistical modeling. Up-to-date details at <https://mc-stan.org/events/stancon2020...> [read more](#)



🔒 📌 Welcome to the Stan Forums!

The Stan Forums provide a community for asking and answering questions about all aspects of Stan. Before creating a new topic please search the Forums to see if your questions has already been answered, or check out the... [read more](#)



Unable to retrieve parameters from a dynamic model

■ Modeling ■ fitting-issues ■ specification



<https://discourse.mc-stan.org/>

## Twitter



**Richard McElreath**  
@rilmcelreath



**Quiche Lorraine,**  
@dan\_p\_simpson



**\mathfrak{Michael}**  
@betanalpha



**EJ Wagenmakers**  
@EJWagenmakers

# Summary

- Computational modeling is never new → don't let it fear you!
- Learn some statistics (e.g., different statistical distributions)
- Learn some math (e.g., linear algebra)
- Learn some programming (e.g., R/Python/Julia/Matlab)
- Learn to seek external help (e.g., existing packages)
- Learn in pairs; practice makes perfect!



**Richard McElreath**  
@rilmcelreath



I say this a lot, bc I am also confused quite often.



**Anna Jacobson** @AnnaChingChing · Feb 21

"If you are confused, it is only because you are trying to understand." -  
@rilmcelreath in Statistical Rethinking



# Acknowledgement



Addaptive  
Learning  
Psychology and  
Neuroscience

## Funding:



Medical  
Research  
Council



Biotechnology and  
Biological Sciences  
Research Council



Economic  
and Social  
Research Council



Centre—UB  
Understanding Human Behaviour



# Thank you!



ANY  
QUESTIONS  
?



Happy Computing!